


NETWORK VIRTUALIZATION

Ethan J. Jackson


THE YEAR 2000

ENTERPRISE DATACENTER



ENTERPRISE DATACENTER

- Static Hosts
- Single Tenant
- Manual configuration
- Networking == Hardware



GOOGLE TRENDS

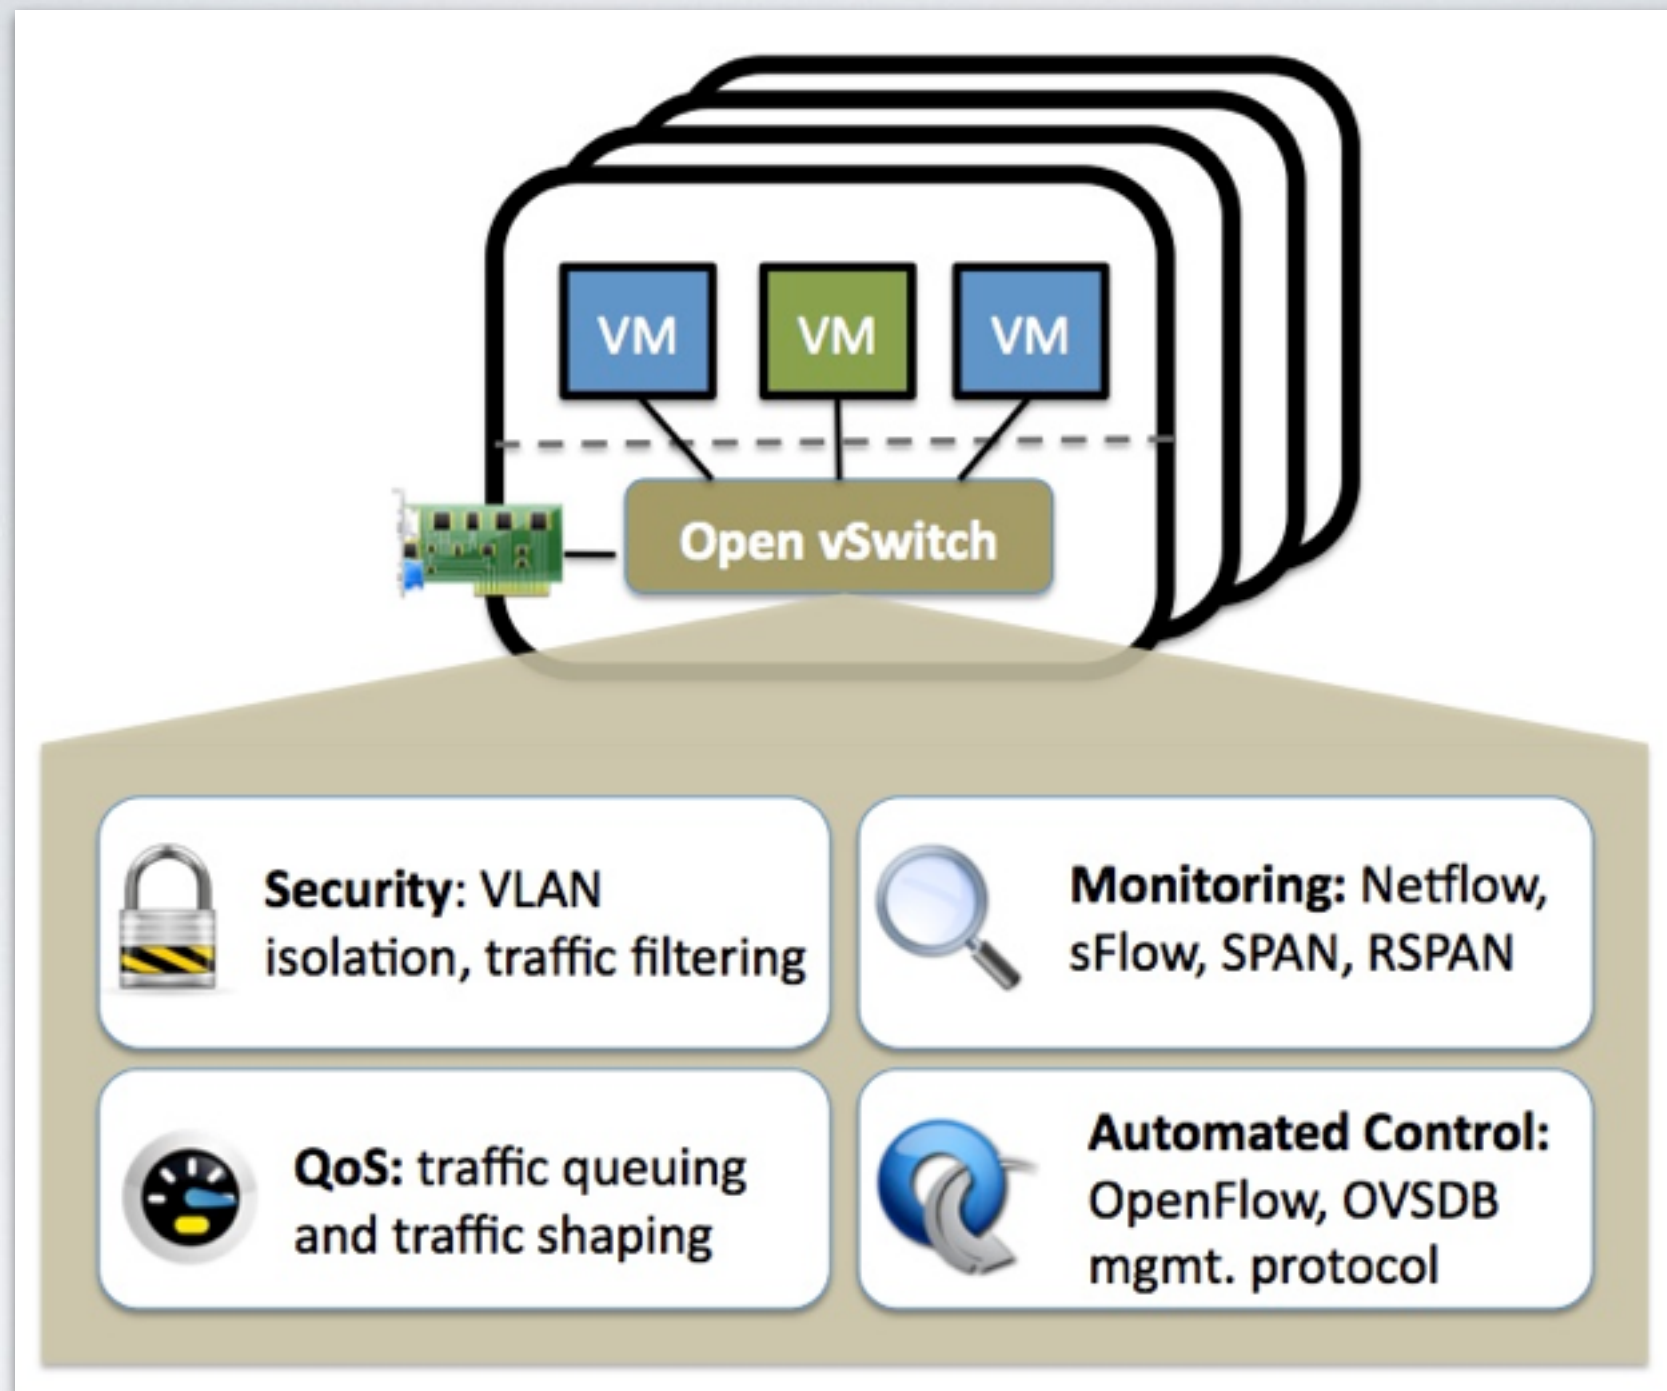
“Hypervisor”

VIRTUALIZED DATACENTER

- Static Hosts <-> Dynamic Hosts
 - Virtual machines boot/die quickly
 - Live migration
 - What does this do to routing?

VIRTUALIZED DATACENTER

- Single Tenant <-> Multi-tenant
 - Public Clouds (Amazon)
 - Private Clouds
 - Unified infrastructure for all applications
- How do we secure this?



VIRTUALIZED DATACENTER

The Hypervisor vSwitch

VIRTUALIZED DATACENTER

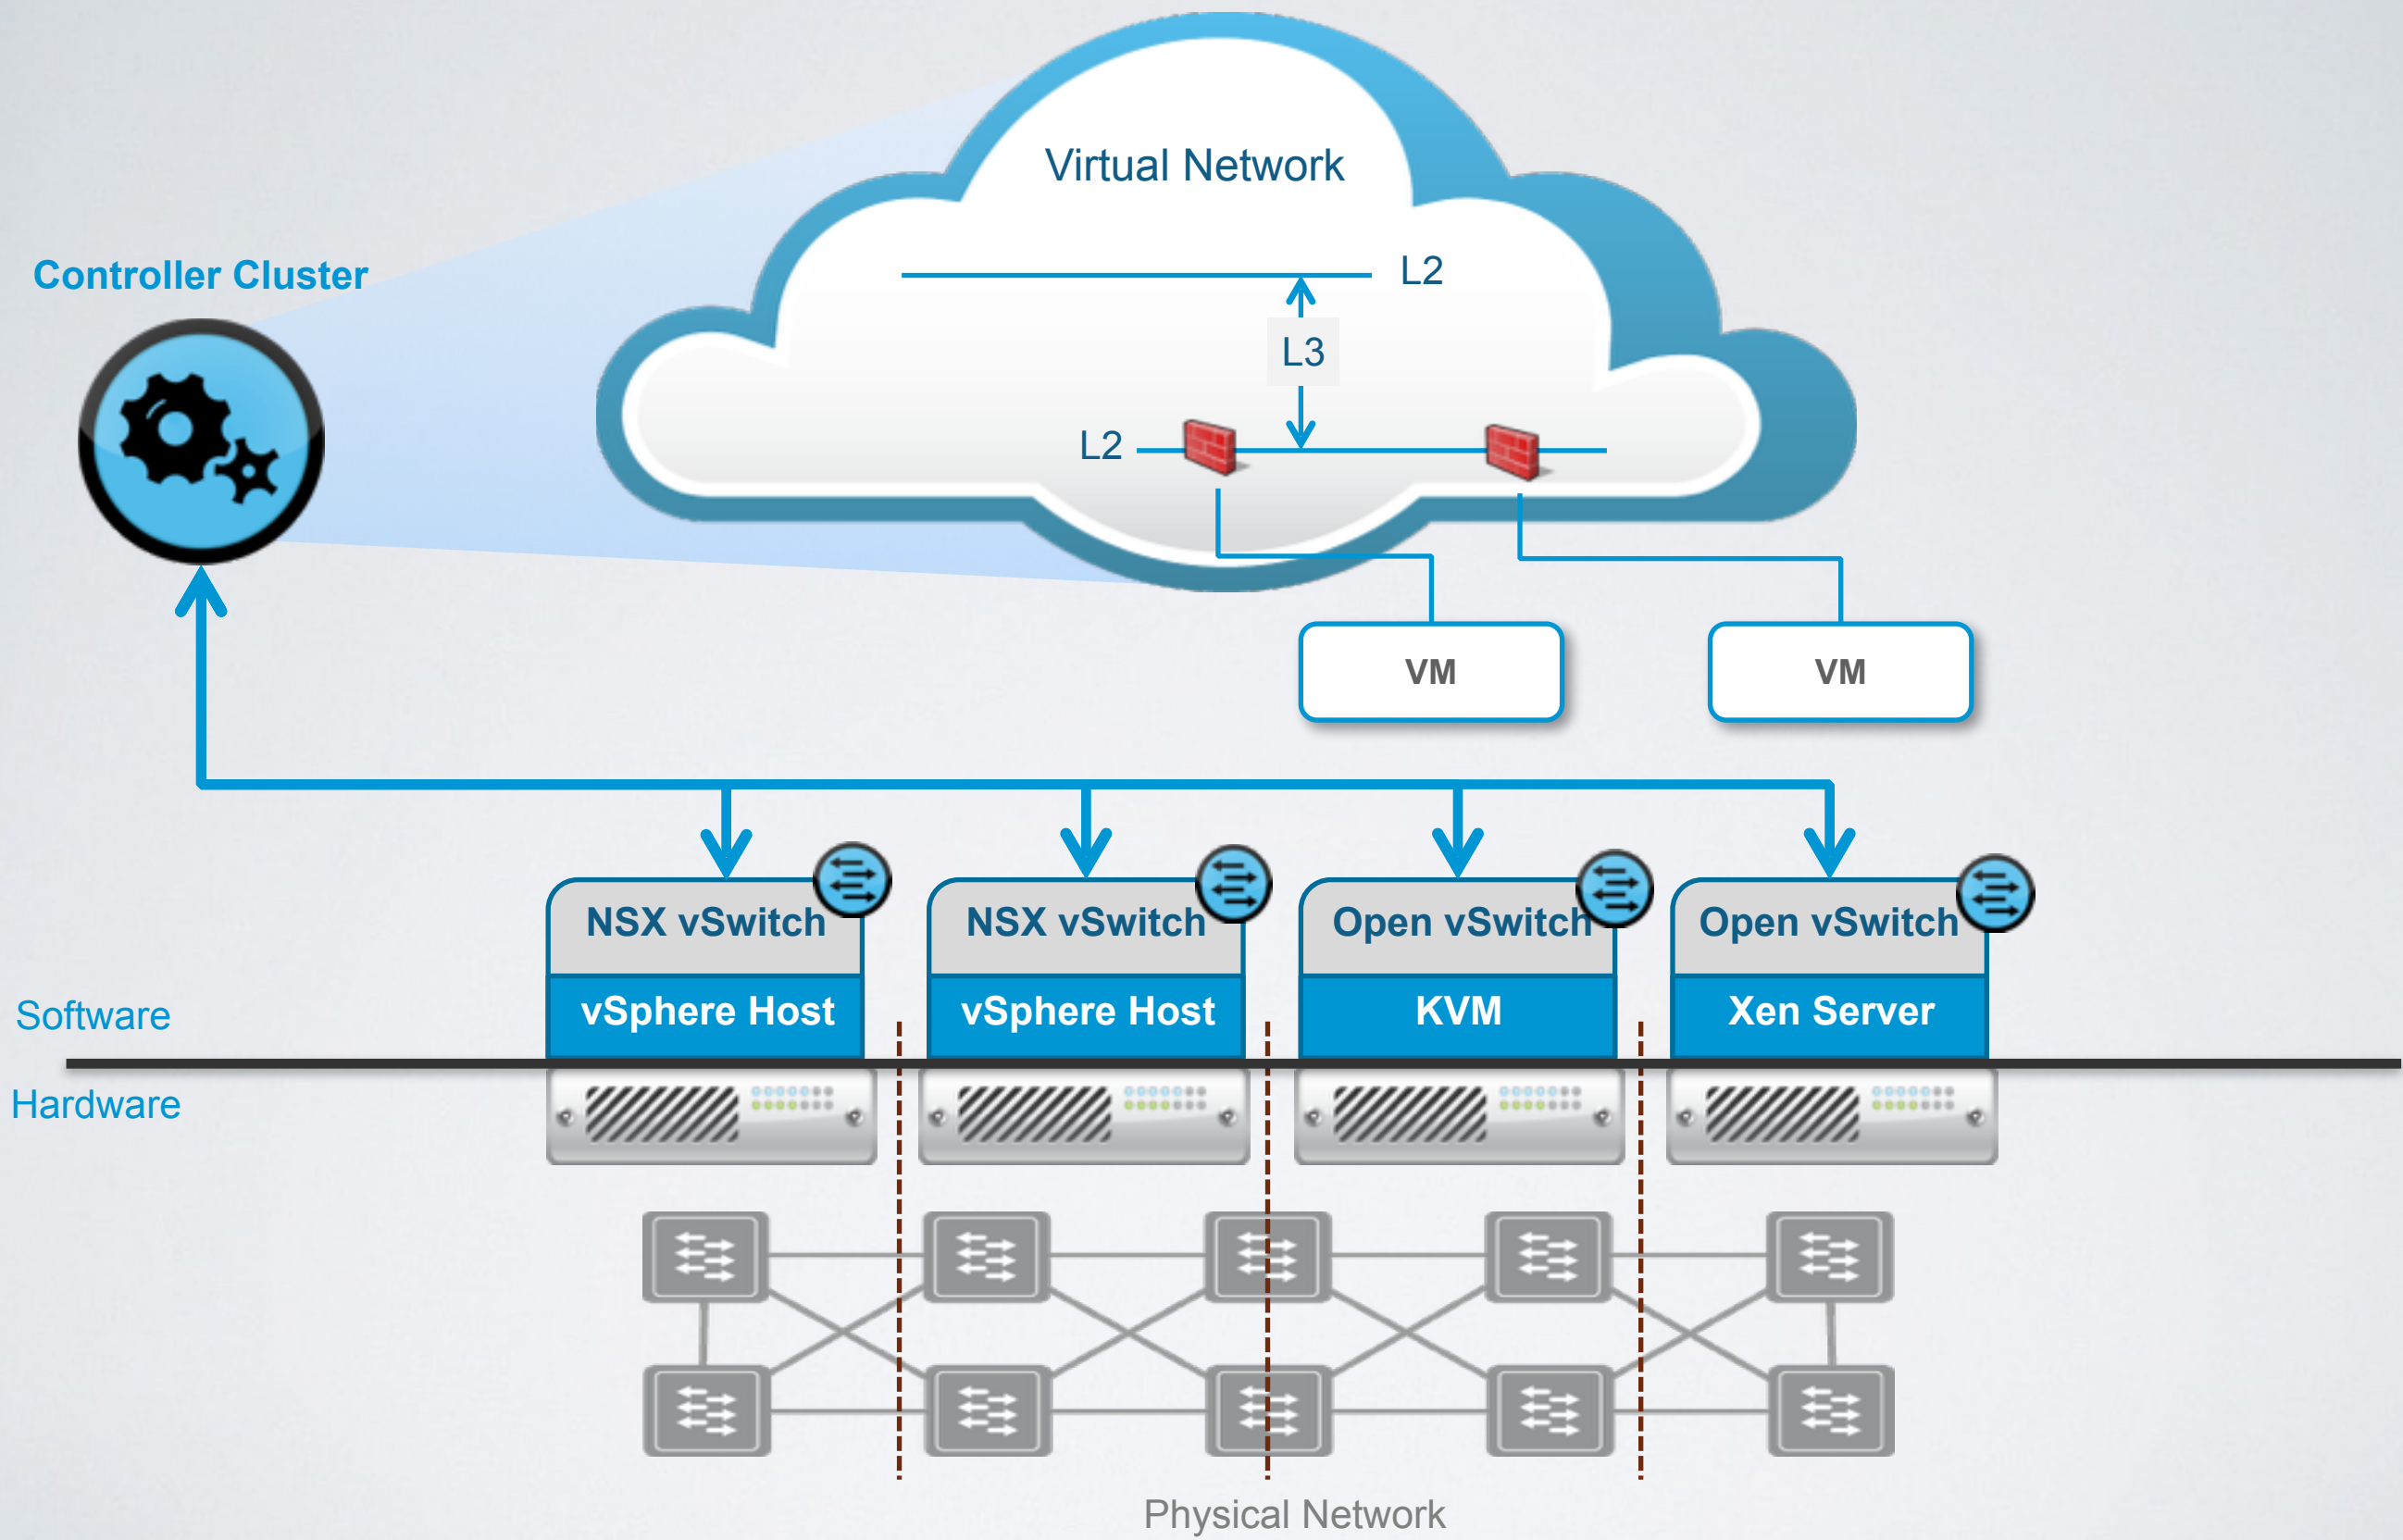
- Networks were broken
 - A high dynamic virtualized cloud
 - Managed manually by network admins

THE PROBLEM

Network Policy is Chained to Network Hardware

THE SOLUTION

Network Virtualization



ABSTRACTION

- Physical network does connectivity (only)
- Virtual network does policy (only)
- Each evolves independently

SIMPLICITY

- Do complex stuff at the edge
 - In software
- Handle global state in a controller
 - Instead of distributed algorithms
- Use cheap, commodity, hardware
 - Instead of Cisco

ADVANTAGES

- Automation (humans cause bugs)
- Deployment Velocity
- Micro-segmentation
- Cloud Bursting
- Inter-DC Migration

THE ARCHITECTURE




THREE PILLARS

- Tunnels
- Virtual Switches
- Controllers


TUNNELS

10.0.0.0/8

10.0.0.0/8




TUNNELS



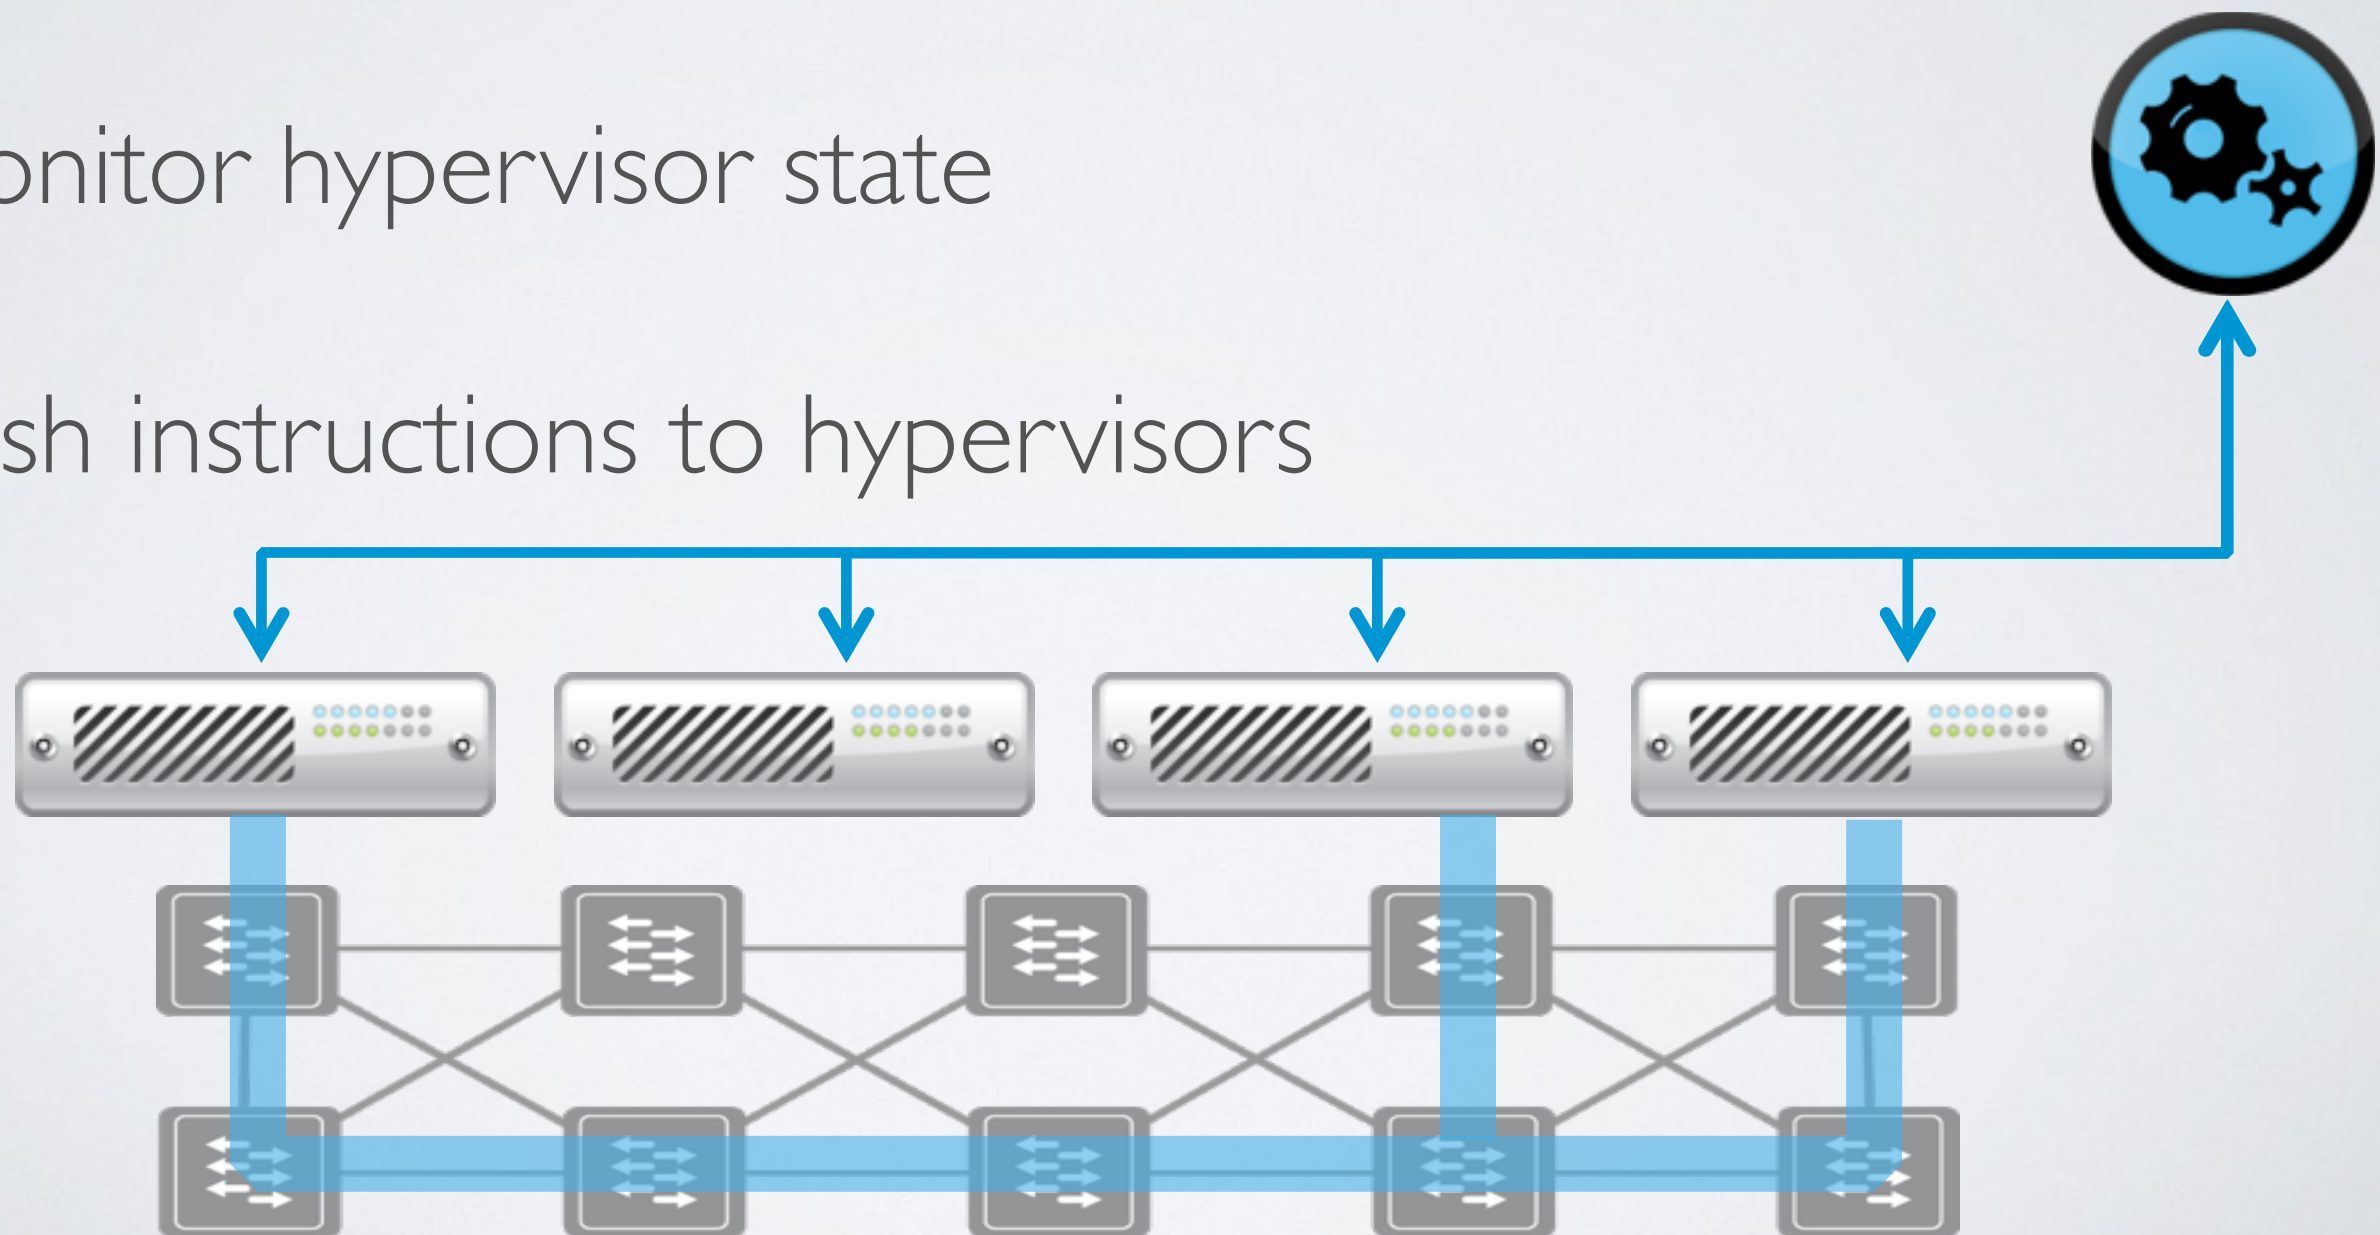
TUNNELS

- Mesh of tunnels
- Which one to use?




CONTROLLERS

- Receive virtual network configuration
- Monitor hypervisor state
- Push instructions to hypervisors




CONTROLLERS



TRADITIONAL vSWITCHES

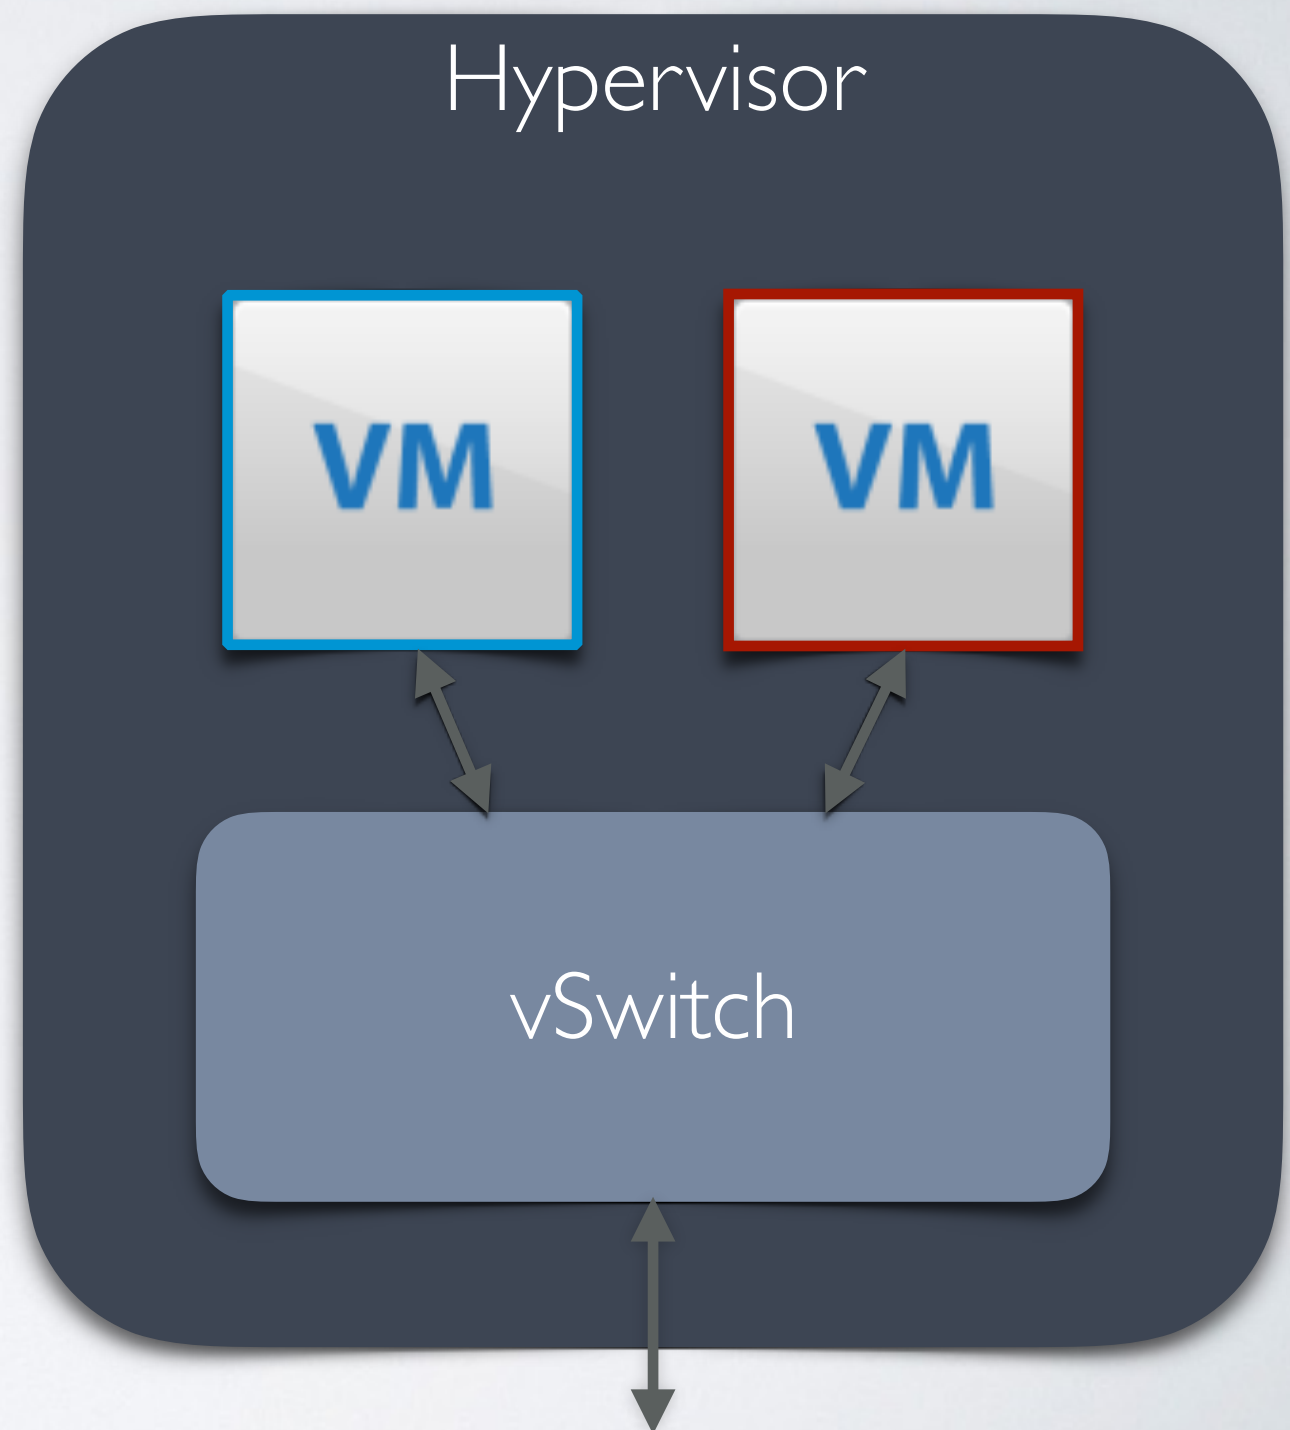
- Core component of hypervisors
- Connect VMs to each other
 - And the physical network
- Traditionally simple L2
 - Linux Bridge



MODERN vSWITCHES

- Implement Controller Policy
 - Create Tunnels
 - Implement Logical Pipeline
- Monitor Remote Hypervisors

Programmable

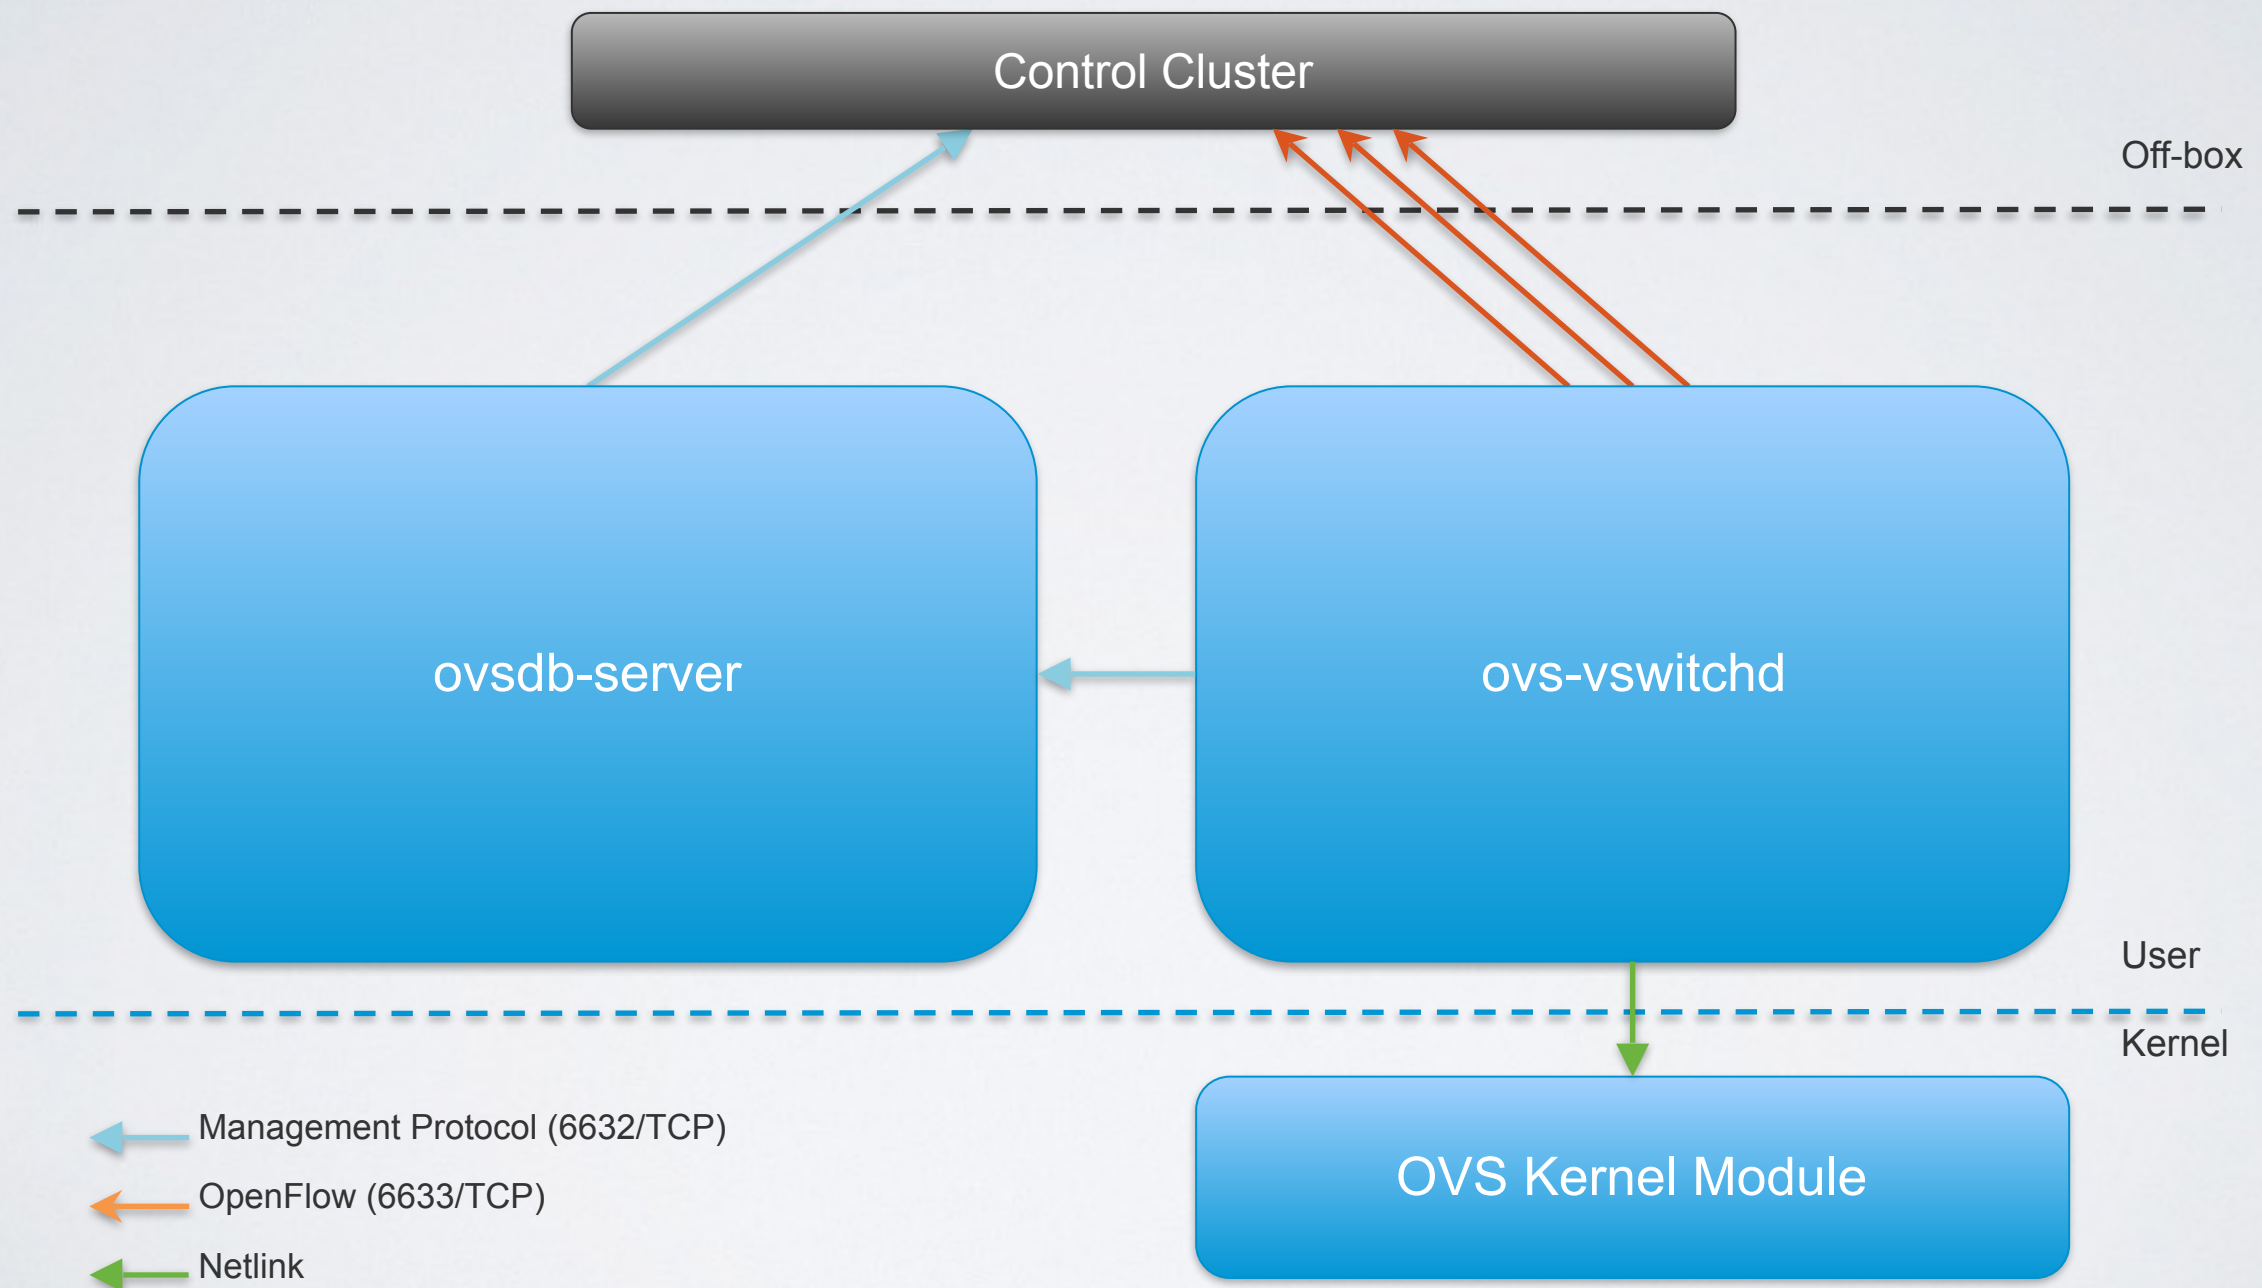



OPEN VSWITCH




- A programmable vSwitch designed for Network Virtualization
 - (And Software Defined Networking)
- Extremely successful
 - Bundled with the Linux Kernel
 - Default switch for KVM deployments

OPEN VSWITCH



OVSDDB



Control Cluster

ovsdb-server

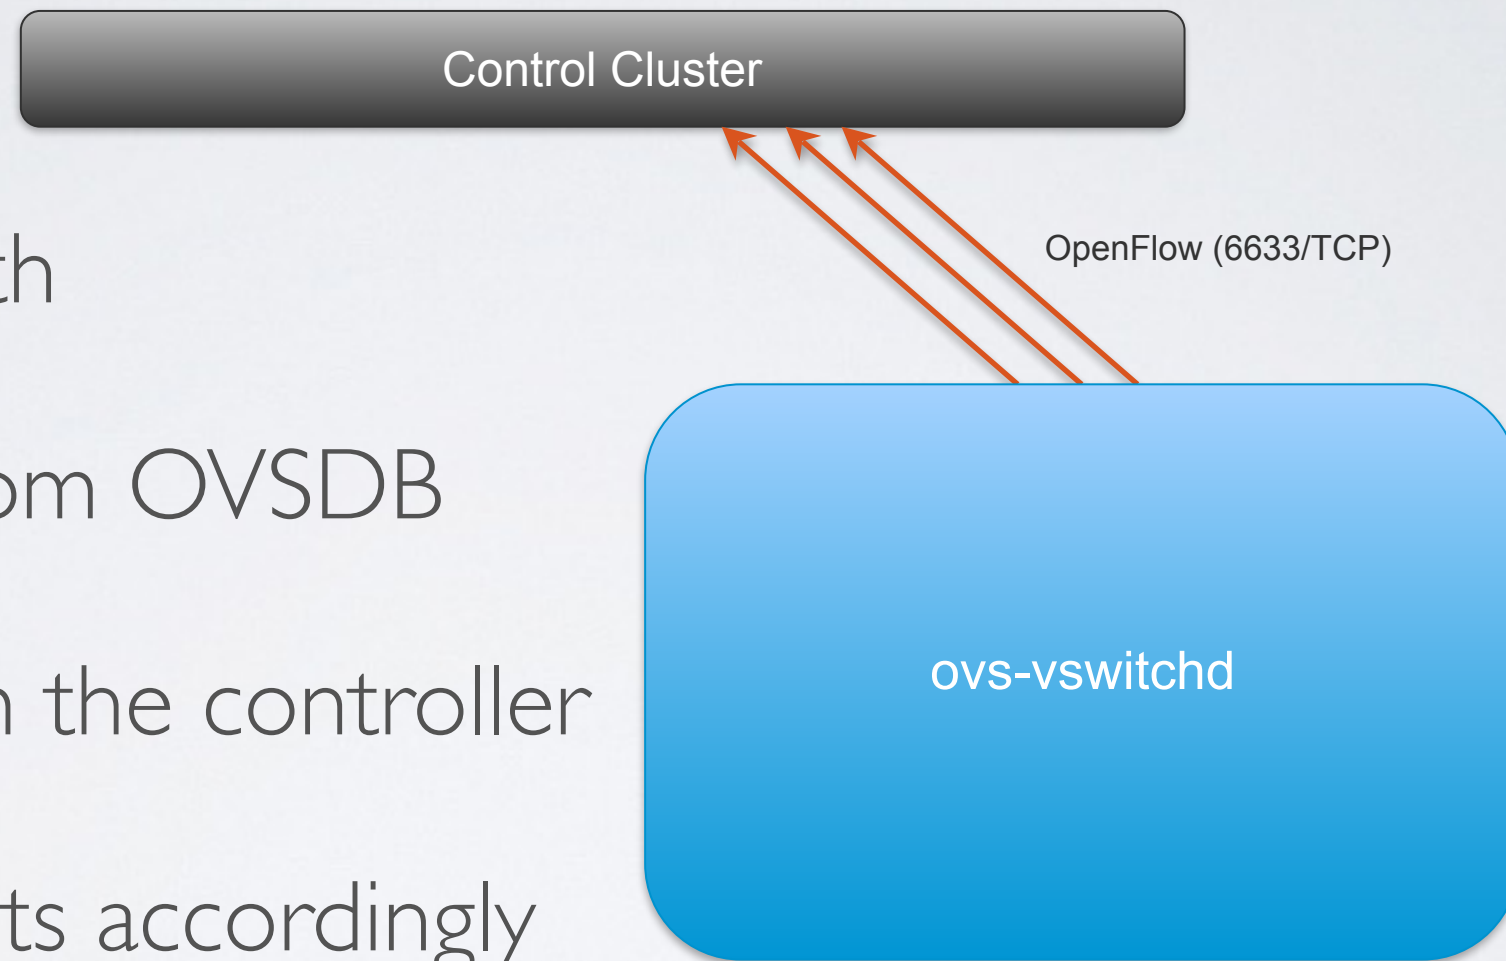
- Transactional database
- Speaks JSON RPC
- Stores slow-changing config
 - Ports
 - Tunnels

← Management Protocol (6632/TCP)

OVS VSWITCHD



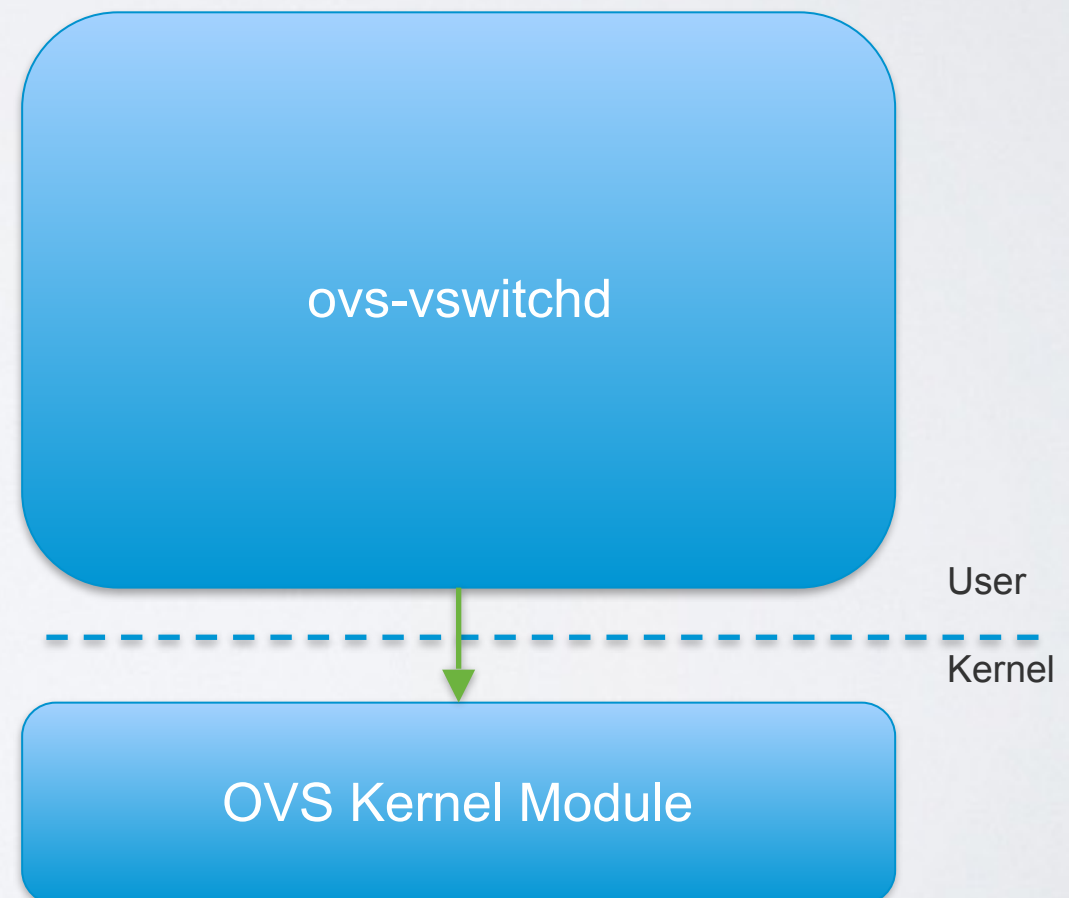
- Switch Slow Path
- Reads config from OVSDB
- OpenFlow from the controller
- Forwards packets accordingly



DATAPATH



- Switch fast path
- Typically implemented in Kernel
- Receives Packets from the NIC and forwards them
- Punts them up if confused



FLOW CACHING




Pipeline

Switch



User

Kernel

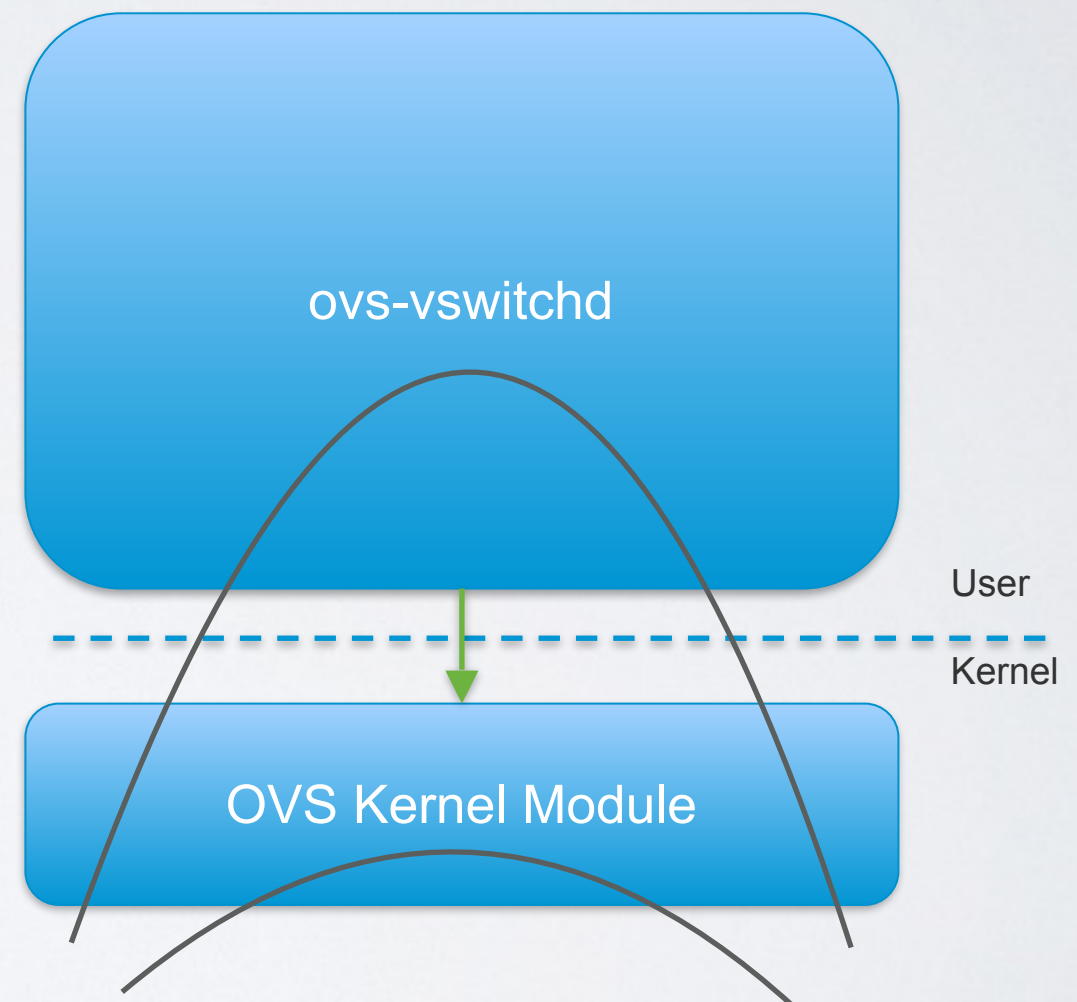


OVS Kernel Module

FLOW CACHING



- Complex processing happens in slow path
- Forwarding happens in datapath
- Balances speed ...
- And programmability



ARCHITECTURE SUMMARY

- Tunnels
 - Allow physical network to forward packets, ignorant of network policy
- Controller
 - Manages global configuration
- vSwitch
 - Highly programmable dataplane does the heavy lifting


THE FUTURE

MIDDLEBOXES

Network Function Virtualization

- Middleboxes (NATS, Firewalls, IDS, IPS, etc)
- Fit Badly into the model (state management is hard)
- Network Function Virtualization
 - Shove middleboxes into virtual machines
 - Use Netmap/DPDK to make it fast

WHITE-BOX SWITCHING



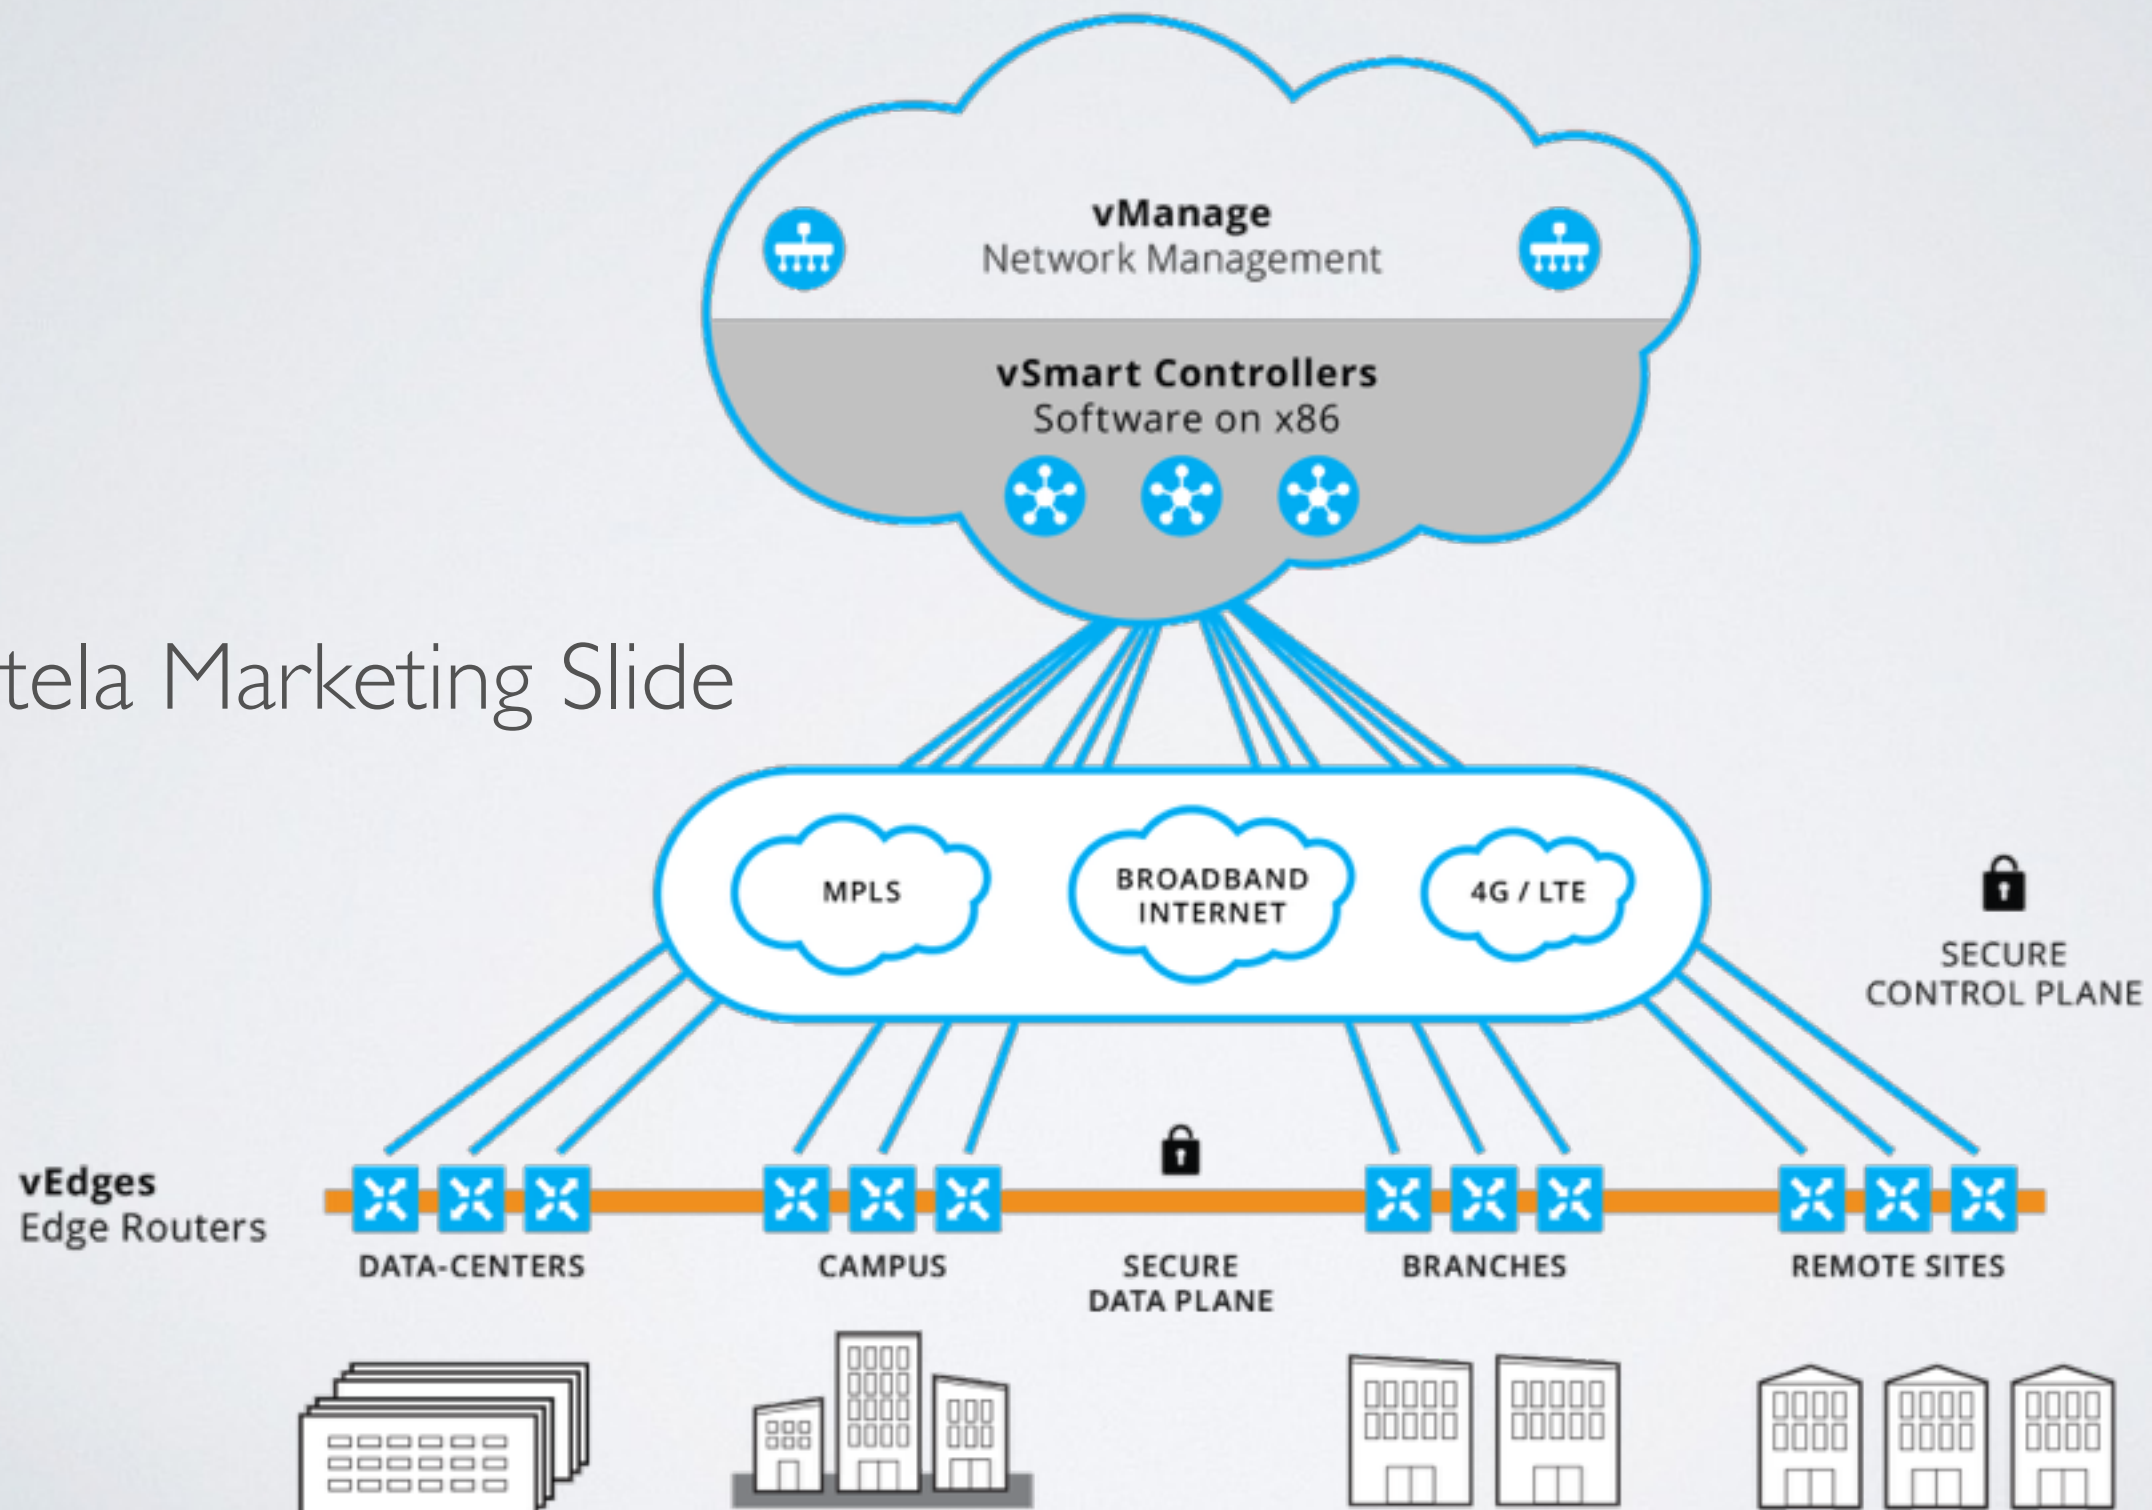
PROGRAMMABLE HARDWARE

- OpenFlow Failed
 - (For hardware)
- ASICs are inflexible
- Enter P4 ...



SD-WAN

Viptela Marketing Slide



WILD SPECULATION

SHAMELESS PLUG

- Found this interesting?
- Good Hacker?
- Find me ...
 - ejj@eecs.berkeley.edu
 - ejj.github.io

THANK YOU!